

Assignment-4

Assignment Date	04 October 2022
Student Name	Alagappan N
Student Roll Number	2127190801005
Maximum Marks	2 Marks

Question-1:

Download the dataset

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Question-2:

Load the dataset

```
In [2]: data = pd.read_csv(r"Mall_Customers.csv")
```

```
In [3]: data.head();
```

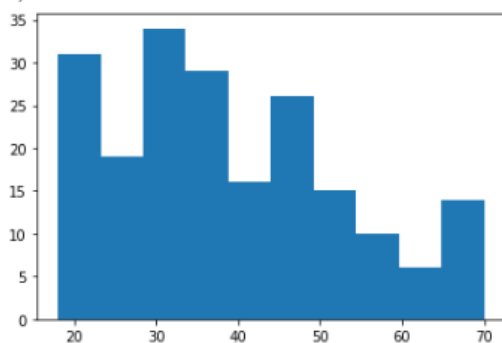
Question 3:

Perform Below Visualizations.

1) Univariate Analysis

```
In [4]: plt.hist(data['Age'])
```

```
Out[4]: (array([31., 19., 34., 29., 16., 26., 15., 10., 6., 14.]),
array([18., 23.2, 28.4, 33.6, 38.8, 44., 49.2, 54.4, 59.6, 64.8, 70. ]),
)
```



2) Bi-variate analysis

```
In [5]: plt.scatter(data['Gender'],data['Annual Income (k$)'])
```

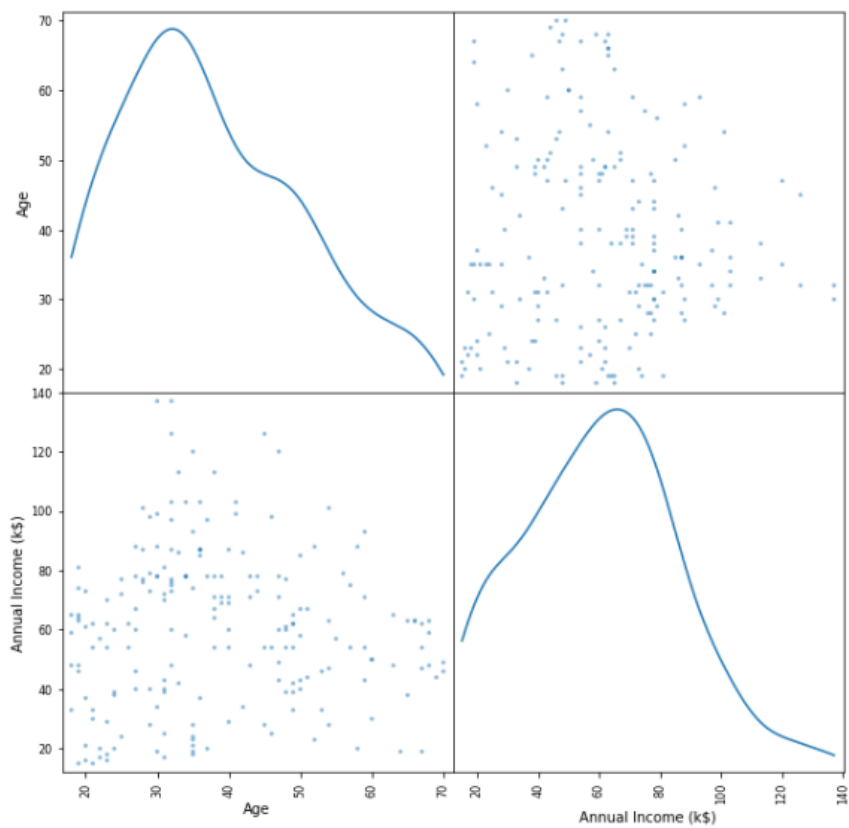
Out[5]:



3) Multi-variate analysis

```
In [6]: pd.plotting.scatter_matrix(data.loc[:, "Age": "Annual Income (k$)" ], diagonal = "kde", figsize=(10,10))
```

Out[6]: array([[
],
],
],
],
 dtype=object)



Question 4:

Perform descriptive statistics on the dataset

```
In [7]: data.describe()
```

```
Out[7]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
In [8]: data.describe().T
```

```
Out[8]:
```

	count	mean	std	min	25%	50%	75%	max
CustomerID	200.0	100.50	57.879185	1.0	50.75	100.5	150.25	200.0
Age	200.0	38.85	13.969007	18.0	28.75	36.0	49.00	70.0
Annual Income (k\$)	200.0	60.56	26.264721	15.0	41.50	61.5	78.00	137.0
Spending Score (1-100)	200.0	50.20	25.823522	1.0	34.75	50.0	73.00	99.0

Question 5:

Check for missing values

and deal with them

```
In [9]: data.isna().sum()
```

```
Out[9]: CustomerID      0
Gender      0
Age      0
Annual Income (k$)      0
Spending Score (1-100)  0
dtype: int64
```

Question 6:

Find the outliers and replace them outliers

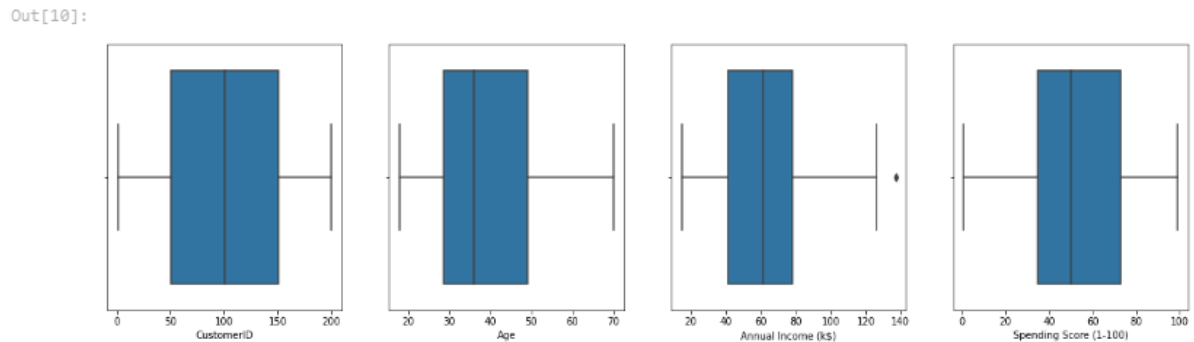
```
In [10]: fig,ax=plt.subplots(figsize=(25,5))

plt.subplot(1, 5, 2)
sns.boxplot(x=data['Age'])

plt.subplot(1, 5, 3)
sns.boxplot(x=data['Annual Income (k$)'])

plt.subplot(1, 5, 4)
sns.boxplot(x=data['Spending Score (1-100)'])

plt.subplot(1, 5, 1)
sns.boxplot(x=data['CustomerID'])
```



Handling Outlier

```
In [11]: quant=data.quantile(q=[0.25,0.75])
quant
```

```
Out[11]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
0.25	50.75	28.75	41.5	34.75
0.75	150.25	49.00	78.0	73.00

```
In [12]: quant.loc[0.75]
```

```
Out[12]: CustomerID      150.25
Age                49.00
Annual Income (k$)   78.00
Spending Score (1-100) 73.00
Name: 0.75, dtype: float64
```

```
In [13]: quant.loc[0.25]
```

```
Out[13]: CustomerID      50.75
Age                28.75
Annual Income (k$)   41.50
Spending Score (1-100) 34.75
Name: 0.25, dtype: float64
```

```
In [14]: iqr=quant.loc[0.75]-quant.loc[0.25]
iqr
```

```
Out[14]: CustomerID      99.50
Age                20.25
Annual Income (k$)   36.50
Spending Score (1-100) 38.25
dtype: float64
```

```
In [15]: low=quant.loc[0.25]-(1.5 *iqr)
low
```

```
Out[15]: CustomerID      -98.500
Age                -1.625
Annual Income (k$)  -13.250
Spending Score (1-100) -22.625
dtype: float64
```

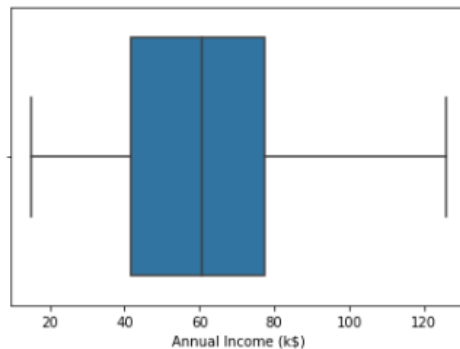
```
In [16]: up=quant.loc[0.75]+(1.5 *iqr)
up
```

```
Out[16]: CustomerID      299.500
Age          79.375
Annual Income (k$)    132.750
Spending Score (1-100) 130.375
dtype: float64
```

```
In [18]: data['Annual Income (k$)']= np.where(data['Annual Income (k$)']>132,60,data['Annual Income (k$)'])
```

```
In [19]: sns.boxplot(x=data['Annual Income (k$)'])
```

Out[19]:



Question 7:

Check for Categorical columns and perform encoding.

```
In [20]: data.info()
```

```
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
In [21]: data['Gender'].unique()
```

```
Out[21]: array(['Male', 'Female'], dtype=object)
```

```
In [22]: data['Gender'].replace({'Male':1,"Female":0},inplace=True)
```

```
In [23]: data
```

```
Out[23]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40
...
195	196	0	35	120	79
196	197	0	45	126	28
197	198	1	32	126	74
198	199	1	32	60	18
199	200	1	30	60	83

200 rows × 5 columns

Question 8:

Scaling the data

```
In [24]: from sklearn.preprocessing import MinMaxScaler  
sc=MinMaxScaler()
```

```
In [25]: df=sc.fit_transform(data.iloc[:,1:])
```

```
In [26]: df
```

```
Out[26]: array([[1.          , 0.01923077, 0.          , 0.3877551 ],  
                [1.          , 0.05769231, 0.          , 0.81632653],  
                [0.          , 0.03846154, 0.00900901, 0.05102041],  
                [0.          , 0.09615385, 0.00900901, 0.7755102 ],  
                [0.          , 0.25          , 0.01801802, 0.39795918],  
                [0.          , 0.07692308, 0.01801802, 0.76530612],  
                [0.          , 0.32692308, 0.02702703, 0.05102041],  
                [0.          , 0.09615385, 0.02702703, 0.94897959],  
                [1.          , 0.88461538, 0.03603604, 0.02040816],  
                [0.          , 0.23076923, 0.03603604, 0.7244898 ],  
                [1.          , 0.94230769, 0.03603604, 0.13265306],  
                [0.          , 0.32692308, 0.03603604, 1.          ],  
                [0.          , 0.76923077, 0.04504505, 0.14285714],  
                [0.          , 0.11538462, 0.04504505, 0.7755102 ],  
                [1.          , 0.36538462, 0.04504505, 0.12244898],  
                [1.          , 0.07692308, 0.04504505, 0.79591837],  
                [0.          , 0.32692308, 0.05405405, 0.34693878],  
                [1.          , 0.03846154, 0.05405405, 0.66326531],  
                [1.          , 0.65384615, 0.07207207, 0.28571429],  
                [0.          , 0.32692308, 0.07207207, 0.98979592],  
                [1.          , 0.32692308, 0.08108108, 0.34693878],  
                [1.          , 0.13461538, 0.08108108, 0.73469388],  
                [0.          , 0.53846154, 0.09009009, 0.04081633],  
                [1.          , 0.25          , 0.09009009, 0.73469388],  
                [0.          , 0.69230769, 0.11711712, 0.13265306],  
                [1.          , 0.21153846, 0.11711712, 0.82653061],  
                [0.          , 0.51923077, 0.11711712, 0.31632653],  
                [1.          , 0.32692308, 0.11711712, 0.6122449 ],  
                [0.          , 0.42307692, 0.12612613, 0.30612245],  
                [0.          , 0.09615385, 0.12612613, 0.87755102],  
                [1.          , 0.80769231, 0.13513514, 0.03061224],  
                [0.          , 0.05769231, 0.13513514, 0.73469388],  
                [1.          , 0.67307692, 0.16216216, 0.03061224],  
                [1.          , 0.          , 0.16216216, 0.92857143],  
                [0.          , 0.59615385, 0.16216216, 0.13265306],  
                [0.          , 0.05769231, 0.16216216, 0.81632653],  
                [0.          , 0.46153846, 0.17117117, 0.16326531],  
                [0.          , 0.23076923, 0.17117117, 0.73469388],  
                [0.          , 0.34615385, 0.1981982 , 0.25510204],  
                [0.          , 0.03846154, 0.1981982 , 0.75510204],  
                [0.          , 0.90384615, 0.20720721, 0.34693878],  
                [1.          , 0.11538462, 0.20720721, 0.92857143],  
                [1.          , 0.57692308, 0.21621622, 0.35714286],  
                [0.          , 0.25          , 0.21621622, 0.6122449 ],  
                [0.          , 0.59615385, 0.21621622, 0.2755102 ],  
                [0.          , 0.11538462, 0.21621622, 0.65306122],  
                [0.          , 0.61538462, 0.22522523, 0.55102041],  
                [0.          , 0.17307692, 0.22522523, 0.46938776],  
                [0.          , 0.21153846, 0.22522523, 0.41836735],  
                [0.          , 0.25          , 0.22522523, 0.41836735],  
                [0.          , 0.59615385, 0.24324324, 0.52040816],  
                [1.          , 0.28846154, 0.24324324, 0.60204082],  
                [0.          , 0.25          , 0.25225225, 0.54081633],  
                [1.          , 0.78846154, 0.25225225, 0.60204082],  
                [0.          , 0.61538462, 0.25225225, 0.44897959],  
                [1.          , 0.55769231, 0.25225225, 0.40816327],  
                [0.          , 0.63461538, 0.26126126, 0.5          ],
```

```

[1.      , 0.13461538, 0.55855856, 0.1122449 ],
[1.      , 0.19230769, 0.55855856, 0.97959184],
[1.      , 0.57692308, 0.55855856, 0.35714286],
[0.      , 0.26923077, 0.55855856, 0.74489796],
[0.      , 0.30769231, 0.56756757, 0.21428571],
[1.      , 0.30769231, 0.56756757, 0.90816327],
[1.      , 0.48076923, 0.56756757, 0.16326531],
[1.      , 0.40384615, 0.56756757, 0.8877551 ],
[0.      , 0.5       , 0.56756757, 0.19387755],
[0.      , 0.38461538, 0.56756757, 0.76530612],
[0.      , 0.55769231, 0.56756757, 0.15306122],
[0.      , 0.17307692, 0.56756757, 0.89795918],
[1.      , 0.36538462, 0.56756757, 0.       ],
[0.      , 0.23076923, 0.56756757, 0.78571429],
[1.      , 0.30769231, 0.56756757, 0.       ],
[0.      , 0.23076923, 0.56756757, 0.73469388],
[0.      , 0.73076923, 0.57657658, 0.34693878],
[0.      , 0.21153846, 0.57657658, 0.83673469],
[1.      , 0.01923077, 0.59459459, 0.04081633],
[0.      , 0.25       , 0.59459459, 0.93877551],
[1.      , 0.61538462, 0.63063063, 0.25510204],
[0.      , 0.34615385, 0.63063063, 0.75510204],
[1.      , 0.46153846, 0.63963964, 0.19387755],
[0.      , 0.28846154, 0.63963964, 0.95918367],
[0.      , 0.34615385, 0.64864865, 0.26530612],
[1.      , 0.26923077, 0.64864865, 0.63265306],
[1.      , 0.42307692, 0.64864865, 0.12244898],
[1.      , 0.19230769, 0.64864865, 0.75510204],
[1.      , 0.34615385, 0.64864865, 0.09183673],
[1.      , 0.34615385, 0.64864865, 0.92857143],
[0.      , 0.65384615, 0.65765766, 0.12244898],
[0.      , 0.23076923, 0.65765766, 0.86734694],
[1.      , 0.76923077, 0.65765766, 0.14285714],
[1.      , 0.17307692, 0.65765766, 0.69387755],
[1.      , 0.78846154, 0.7027027 , 0.13265306],
[1.      , 0.32692308, 0.7027027 , 0.90816327],
[0.      , 0.36538462, 0.73873874, 0.31632653],
[0.      , 0.26923077, 0.73873874, 0.86734694],
[1.      , 0.53846154, 0.74774775, 0.14285714],
[0.      , 0.21153846, 0.74774775, 0.8877551 ],
[0.      , 0.44230769, 0.75675676, 0.3877551 ],
[1.      , 0.23076923, 0.75675676, 0.97959184],
[0.      , 0.69230769, 0.77477477, 0.23469388],
[1.      , 0.19230769, 0.77477477, 0.68367347],
[0.      , 0.44230769, 0.79279279, 0.16326531],
[0.      , 0.34615385, 0.79279279, 0.85714286],
[0.      , 0.30769231, 0.79279279, 0.2244898 ],
[0.      , 0.26923077, 0.79279279, 0.69387755],
[1.      , 0.28846154, 0.88288288, 0.07142857],
[0.      , 0.38461538, 0.88288288, 0.91836735],
[0.      , 0.55769231, 0.94594595, 0.15306122],
[0.      , 0.32692308, 0.94594595, 0.79591837],
[0.      , 0.51923077, 1.       , 0.2755102 ],
[1.      , 0.26923077, 1.       , 0.74489796],
[1.      , 0.26923077, 0.40540541, 0.17346939],
[1.      , 0.23076923, 0.40540541, 0.83673469]]))

```

Question 9:

Perform any of the clustering algorithms

Kmeans_clustering

```
In [27]: from sklearn.cluster import KMeans
```

```
In [28]: TWSS=[]
k=list(range(2,9))

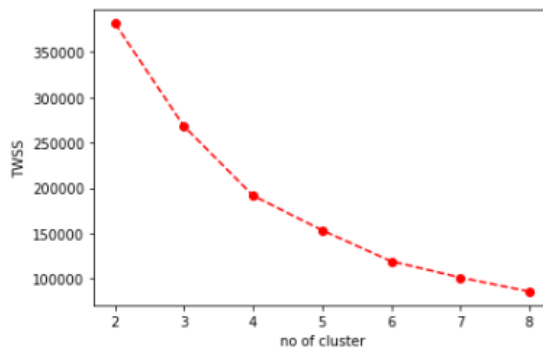
for i in k:
    kmeans=KMeans(n_clusters=i,init='k-means++')
    kmeans.fit(data)
    TWSS.append(kmeans.inertia_)
```

In [29]: TWSS

```
Out[29]: [381550.6840684068,
268082.56760639744,
191612.56821803437,
153394.66603206735,
119223.63779954854,
101364.2432178932,
85819.89345888031]
```

```
In [30]: plt.plot(k,TWSS,'ro--')
plt.xlabel('no of cluster')
plt.ylabel('TWSS')
```

```
Out[30]: Text(0, 0.5, 'TWSS')
```



```
In [31]: #selecting 4 clusters
          model=KMeans(n_clusters=4)
          model.fit(data)
```

```
Out[31]: KMeans(n_clusters=4)
```

```
In [32]: model.labels_
```

[illegible]

```
In [33]: mb=pd.Series(model.labels )
```

```
In [34]: data.head(3)
```

```
Out[34]: CustomerID  Gender  Age  Annual Income (k$)  Spending Score (1-100)
```

0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6

Question 10:

Add the cluster data with the primary dataset

```
In [35]: data['clust']=mb
```

```
In [36]: data.head()
```

```
Out[36]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	clust
0	1	1	19	15	39	1
1	2	1	21	15	81	1
2	3	0	20	16	6	1
3	4	0	23	16	77	1
4	5	0	31	17	40	1

```
In [37]: data.tail()
```

```
Out[37]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	clust
195	196	0	35	120	79	3
196	197	0	45	126	28	2
197	198	1	32	126	74	3
198	199	1	32	60	18	2
199	200	1	30	60	83	3

Question 11:

Split the data into dependent and independent variables.

```
In [38]: #dependent  
y= data['clust']  
y
```

```
Out[38]:
```

0	1
1	1
2	1
3	1
4	1
..	
195	3
196	2
197	3
198	2
199	3

Name: clust, Length: 200, dtype: int32

```
In [39]: #independent  
x= data.drop(columns=['CustomerID','clust'],axis=1)  
x.head()
```

```
Out[39]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	19	15	39
1	1	21	15	81
2	0	20	16	6
3	0	23	16	77
4	0	31	17	40

```
In [52]: x.tail()
```

```
Out[52]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
195	0	35	120	79
196	0	45	126	28
197	1	32	126	74
198	1	32	60	18
199	1	30	60	83

Question 12:

Split the model into training and testing

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

Question 13:

Build the model.

```
In [41]: from sklearn.ensemble import RandomForestClassifier
```

```
In [42]: rf=RandomForestClassifier()
```

Question 14:

Train the model.

```
In [117]: rf.fit(x_train,y_train)
```

```
Out[117]: RandomForestClassifier()
```

Question 15:

Test the model.

```
In [118]: #prediction  
pred=rf.predict(x_test)
```

Question 16:

Measure the performance using Metrics.

```
In [119]: # Accuracy of DI model  
from sklearn.metrics import accuracy_score  
accuracy_score(y_test,pred)
```

```
Out[119]: 0.975
```

```
In [120]: #confusion matrix  
from sklearn import metrics  
metrics.confusion_matrix(y_test,pred)
```

```
Out[120]: array([[13,  0,  0,  0],  
                [ 0, 10,  0,  0],  
                [ 1,  0,  8,  0],  
                [ 0,  0,  0,  8]], dtype=int64)
```